

Deep Learning Workshop


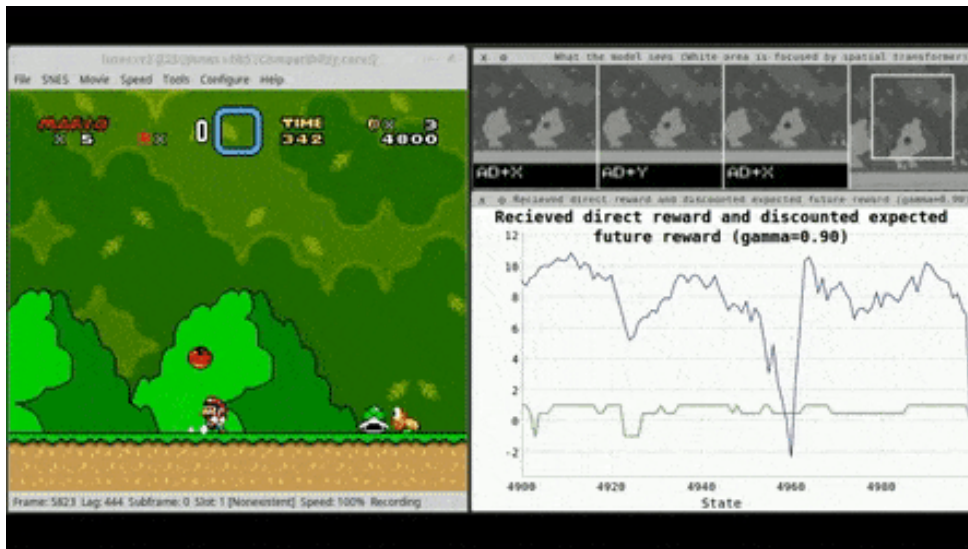
Reinforcement Learning




Instructor: Aaron Low

HELP University, Faculty of Computing and Digital Technology

Reinforcement Learning Example: Video Games



Reinforcement Learning Example: Robotics



Reinforcement Learning Example: Robotics




Not actually reinforcement learning

Reinforcement Learning Success in Games




DeepMind's AlphaGo beats Lee Sedol

Reinforcement Learning Success in Games

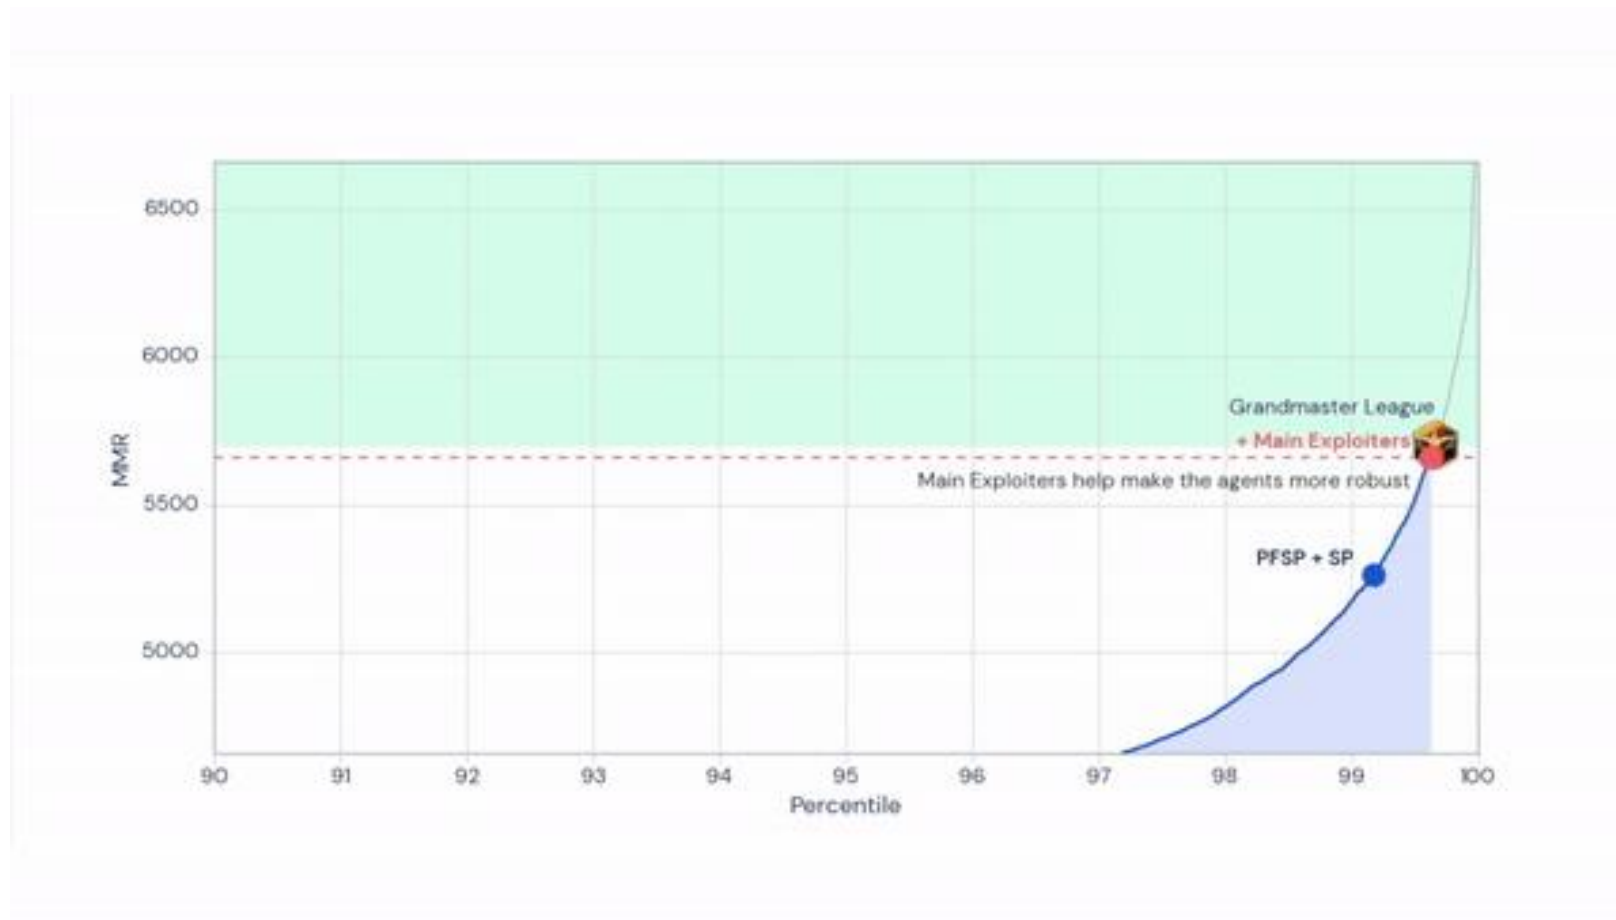


AlphaZero



DeepMind's AlphaZero in Chess, Shogi, Go

Reinforcement Learning Success in Games




DeepMind's AlphaStar in Starcraft 2

Reinforcement Learning Success in Games



OpenAI's OpenAI Five in Dota 2

Reinforcement Learning Success in Games



Facebook and CMU's Pluribus in Poker

Classes of Learning Problems

Supervised Learning

Data: (x, y)
 x is data, y is label

Goal: Learn function to map
 $x \rightarrow y$

Apple example:




This thing is an apple.

Unsupervised Learning

Data: x
 x is data, no labels!

Goal: Learn underlying
structure

Apple example:



This thing is like
the other thing.

Reinforcement Learning

Data: state-action pairs

Goal: Maximize future rewards
over many time steps

Apple example:



Eat this thing because it
will keep you alive.

Reinforcement Learning



Agent and Environment


Agent: The intelligent unit that learns what actions to take

Environment: The surrounding world which provides feedback to the agent

Action: An action the agent takes in the environment e.g. moving


State: The configuration the agent is currently in

Reward: How much benefit taking an action from the current state provides




Deterministic vs Stochastic

Deterministic Grid World




Stochastic Grid World




Markov Decision Process

- Transition function, $T(s'|s, a)$
 - Probability that **action, a** from **state, s** will lead to **next state, s'**
 - Known as the **`model`**



Reinforcement Learning

- In reinforcement learning problems, we **don't** know the **Transitions** and the **Rewards**



Reinforcement Learning

- Learn to **maximize rewards**

Total **Reward**: $R_t = r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^t r_t$

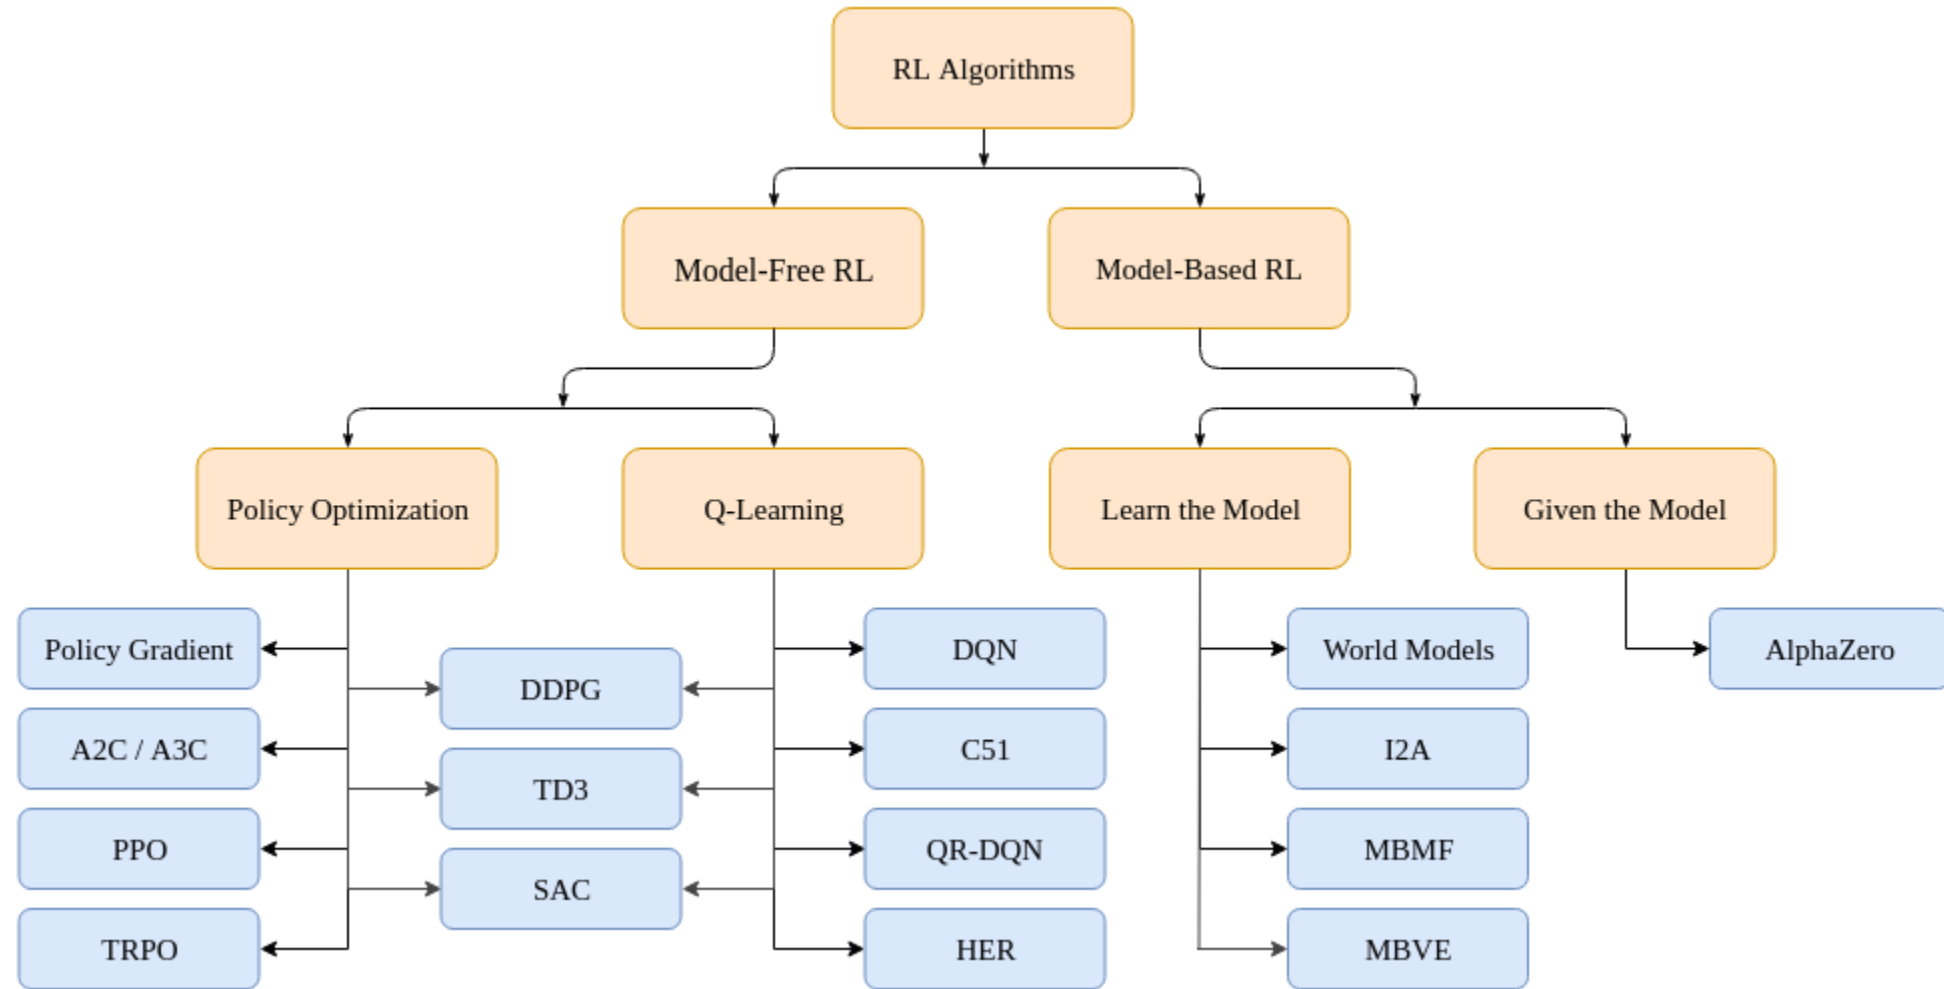
- **Q function**: $Q(s_t, a_t) = E[R_t | s_t, a_t]$
 - Expected total future **reward** an agent in **state**, s_t , receives by making **action**, a_t
- **Policy**: $\pi(s_t) = \operatorname{argmax}_a Q(s_t, a)$
 - The **policy** chooses the best **action** that maximizes future **rewards**
- **Discount factor**: $\gamma, 0 < \gamma < 1$
 - Discount future rewards
 - Immediate rewards are weighted more highly

Importance of Reward: Unintended Consequences



1: MIT Deep Learning Basics: Introduction and Overview Lex Fridman <https://www.youtube.com/watch?v=O5xeyoRL95U>

Reinforcement Learning Taxonomy



Reinforcement Learning Taxonomy

- **Model-based**
 - Learn an approximate model based on experiences
 - Use approximate model to make decisions
 - “I have an idea of where I will be and what the reward will be if I take this action from this state”
- **Model-free**
 - Learn optimal policy or optimal Q-values directly for each action in each state from experience
 - “I know that since I’m in this state, taking this action will be the best”
- Can the agent make predictions about what the next state and reward will be before it takes each action?
 - If yes, then it is **Model-based**

Q-Learning

- Estimate $Q(s, a)$ that **maximizes future reward**
- Use any policy and keep updating (s, a) pairs

How to update Q value

Discount $0 < \gamma < 1$

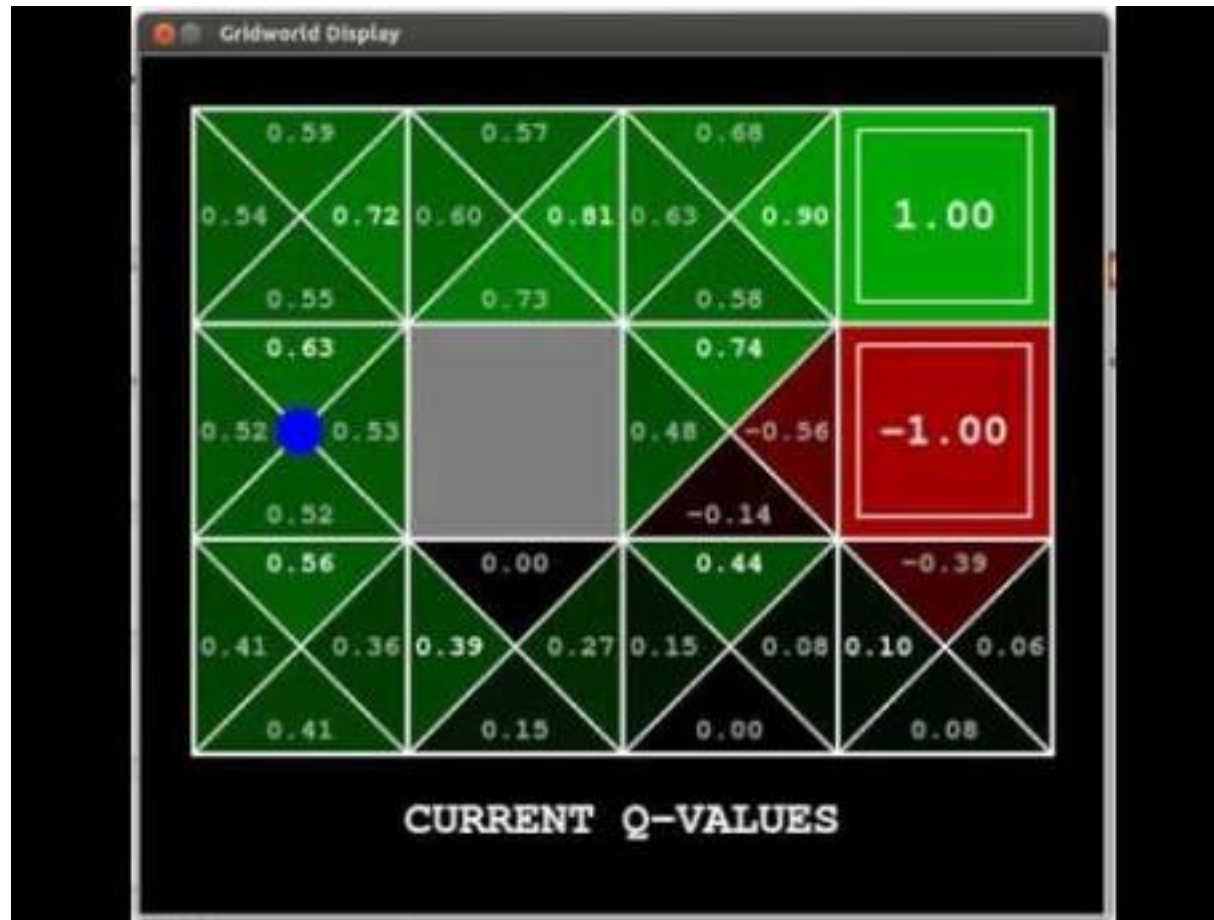
Learning Rate $0 < \alpha < 1$

$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha(R_{t+1} + \gamma \max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t))$$

New Q value Old Q value Reward

The diagram illustrates the Q-learning update equation with several annotations. At the top, 'Discount 0 < γ < 1' has a blue arrow pointing to the γ term in the equation. Below it, 'Learning Rate 0 < α < 1' has a blue arrow pointing to the α term. The equation itself is: Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + α(R_{t+1} + γ max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t)). At the bottom, three blue arrows point upwards to specific parts of the equation: one to Q_{t+1}(s_t, a_t) labeled 'New Q value', one to Q_t(s_t, a_t) labeled 'Old Q value', and one to R_{t+1} labeled 'Reward'.

Q-Learning: Value Iteration




Q-Learning: Value Iteration

- Value Iteration has some **weaknesses**
 - Limited states/actions
 - Doesn't generalize to unseen states
 - **Breakout** game
 - State: screen pixels
 - Image size: **84x84**
 - **4** frames
 - Grayscale (possible levels)
- => **256**^{84*84*4} = 10⁶⁹⁹⁷⁰ >> 10⁸² atoms in the universe



Reinforcement Learning in the Real World

- Training in the real world is not always feasible
- Large number of possible states (we cannot visit all of them and learn about them all)
- Running learning tasks to termination state is NOT always feasible
- Can use simulation to train first before deploying in real world
 - Limited to how well the simulation environment models the real world




Exploration vs Exploitation

- **Exploitation**

- Take deterministic best paths greedily
- At the start (before training), this won't work well as the agent knows nothing about the environment

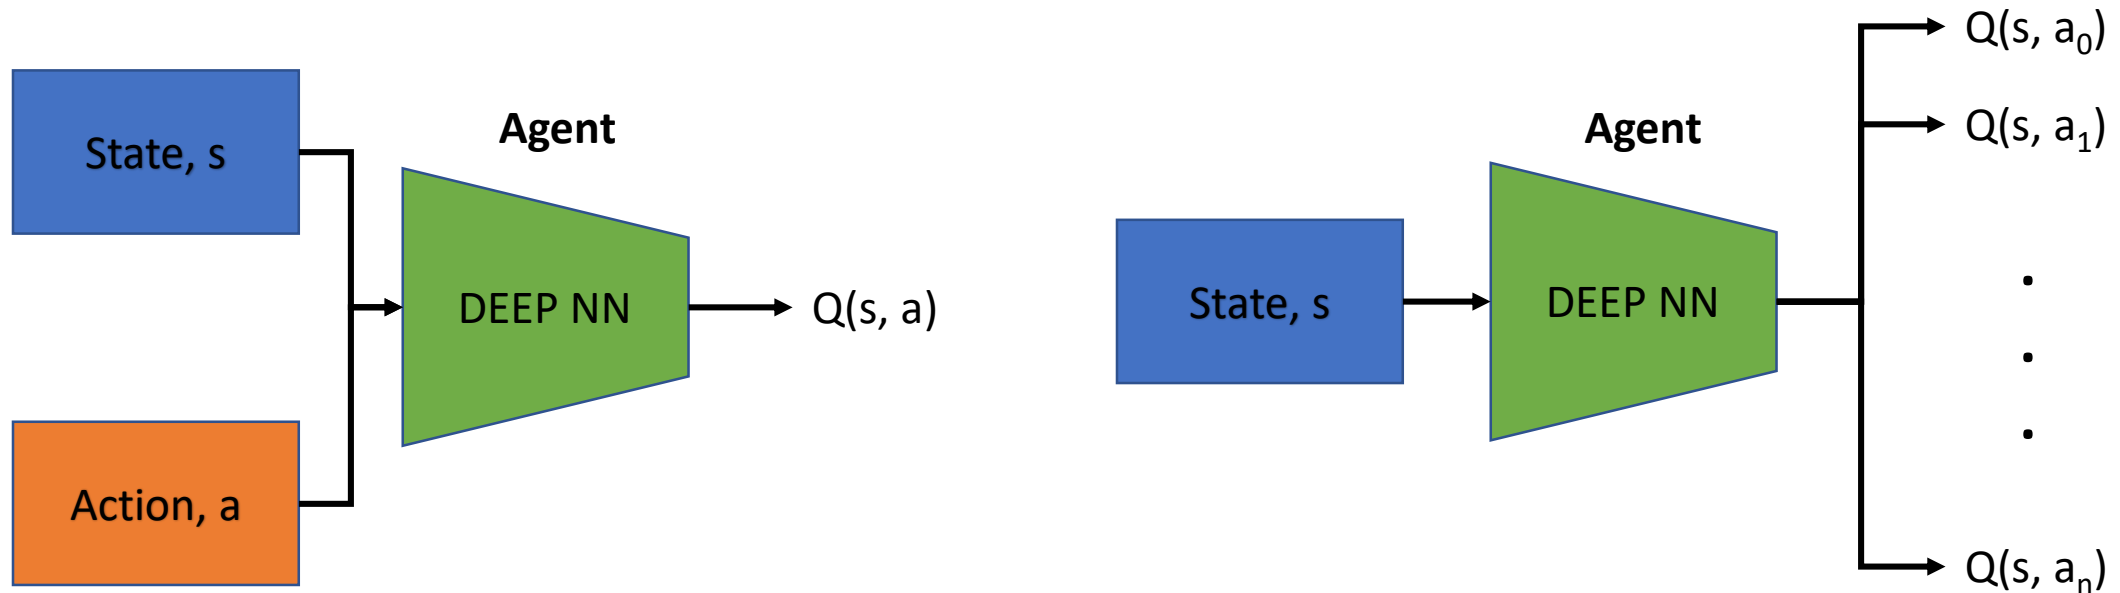
- **Exploration**

- Explore areas we don't know anything about
- Explore by taking random actions at a certain probability, p
- Eventually stop exploring (lower p to 0)



Deep Reinforcement Learning: Deep Q Network (DQN)

Deep Reinforcement Learning: Reinforcement Learning + Neural Networks

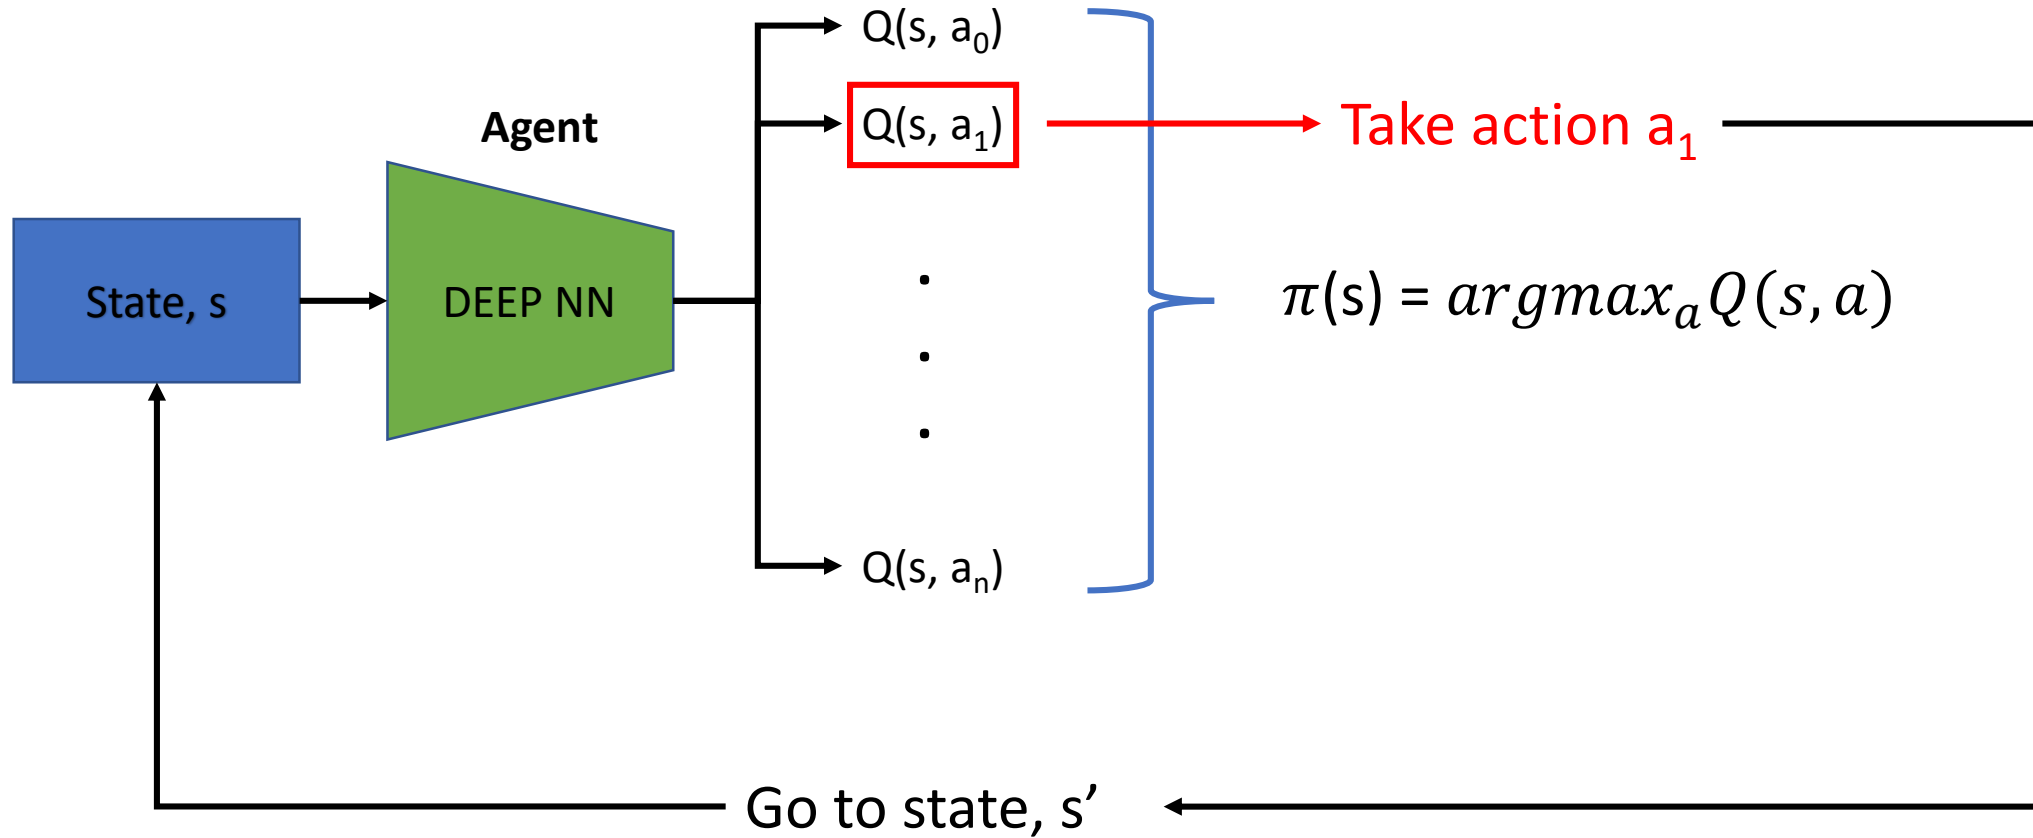


- Use NN to learn Q-function and then use to infer the optimal policy
- Obtain **target** by running agent in environment

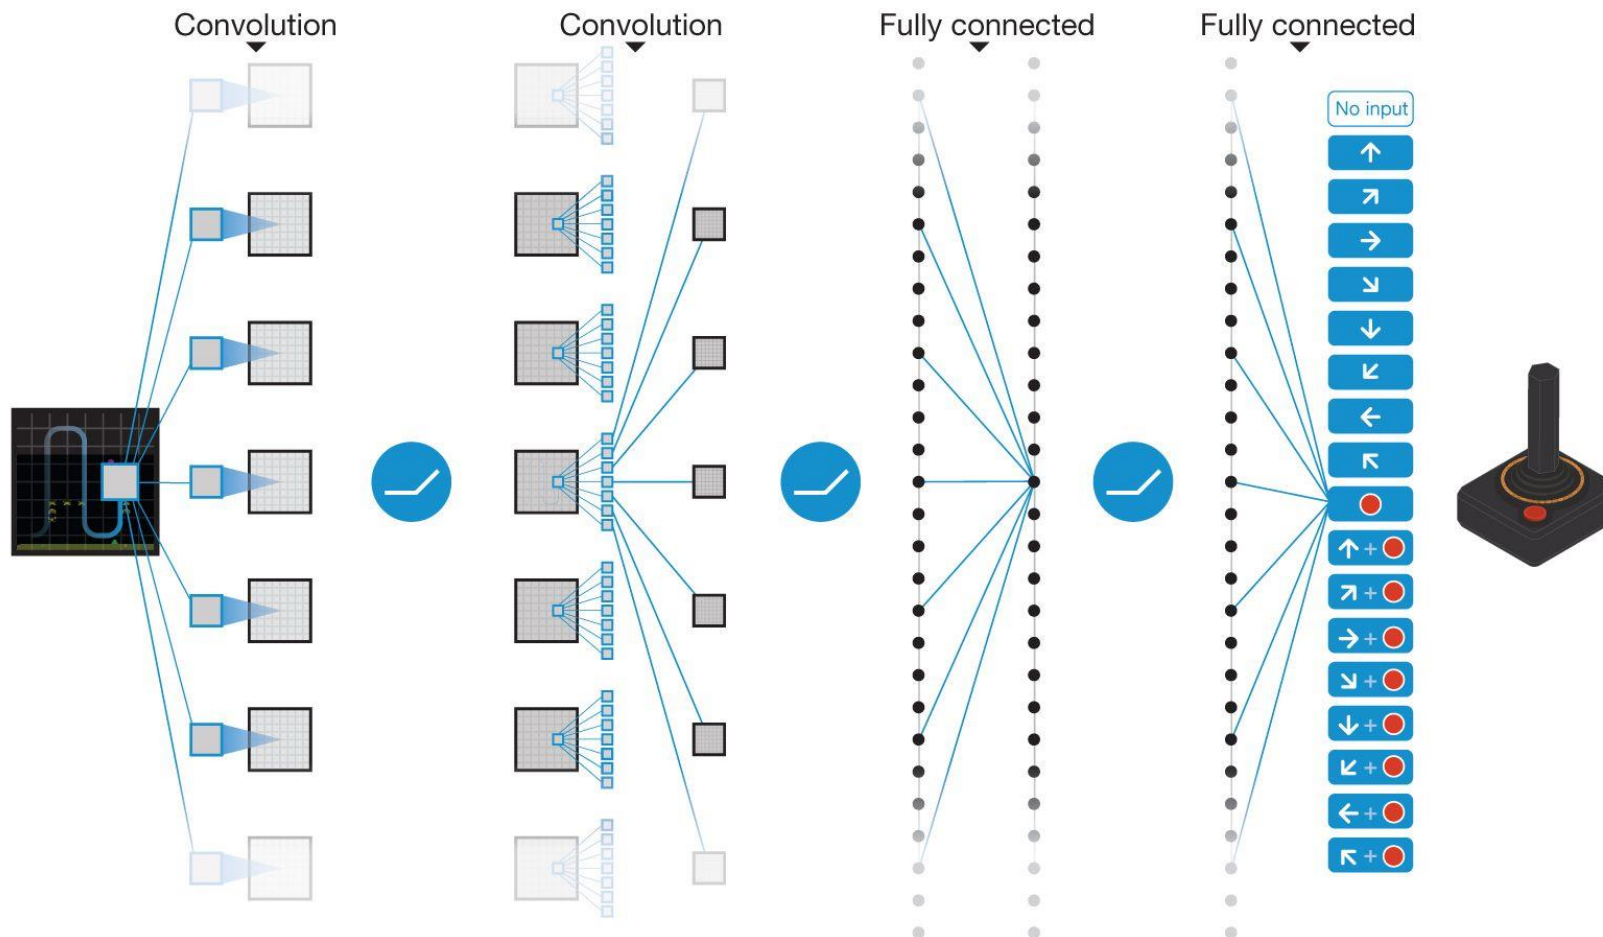
$$\text{Loss Function: } L = E[(\text{target } r + \gamma \max_{a'} Q(s', a') - \text{predicted } Q(s, a))^2]$$

Deep Q Network (DQN)

Example: $Q(s, a_1)$ has highest value




Deep Q Network (DQN)



Deep Q Network (DQN)

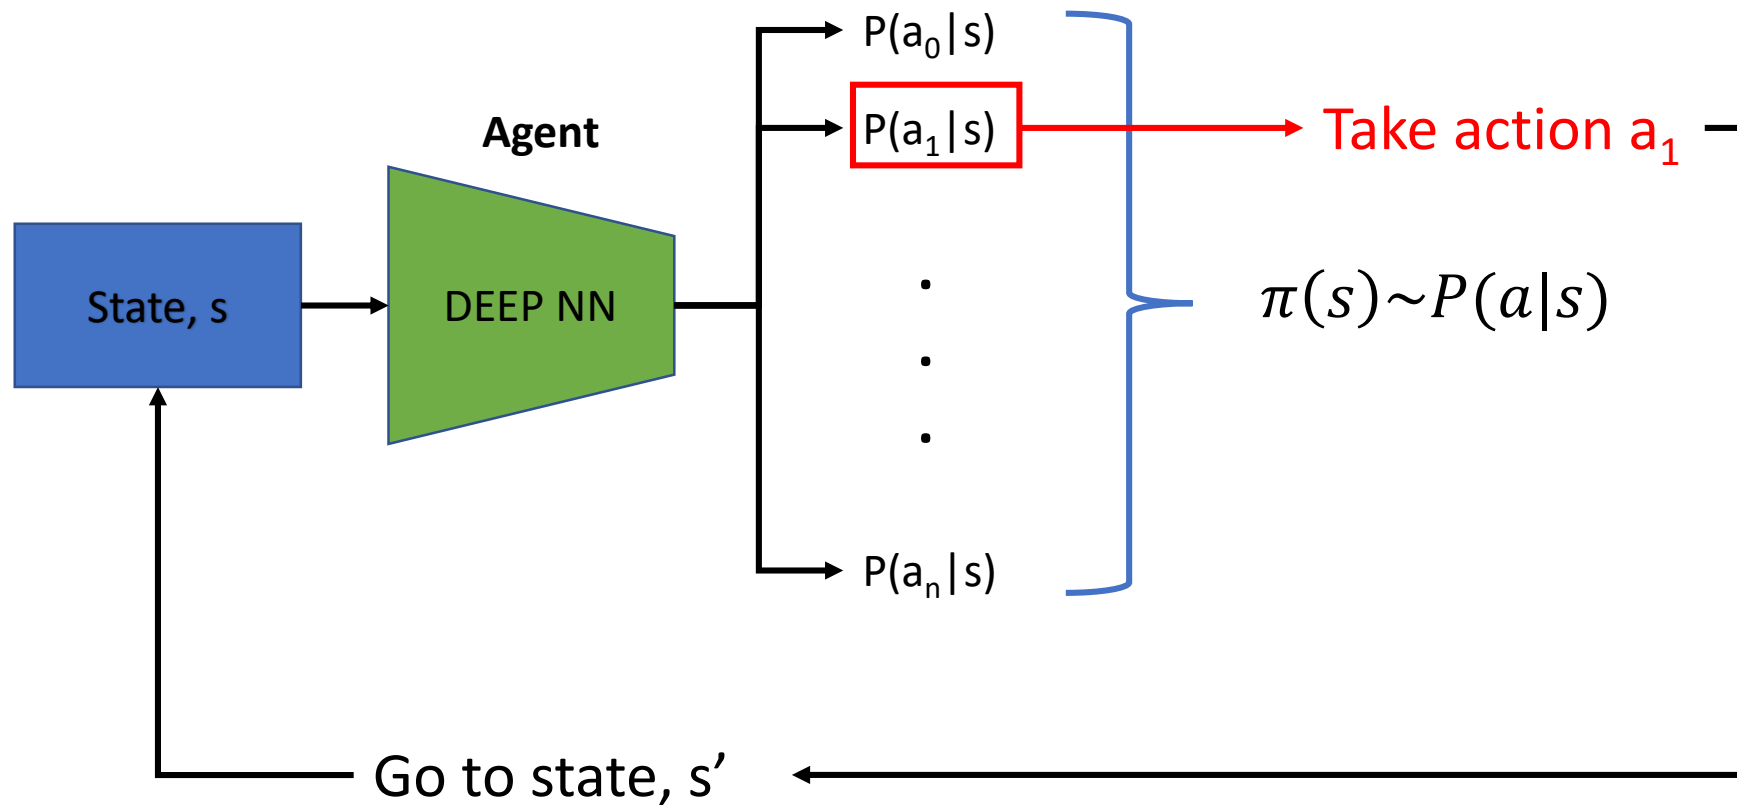
- **Weaknesses**
 - Cannot handle **large** action space
 - Cannot handle **continuous** action space
 - Cannot learn **stochastic** policy (policy is deterministic)



Policy Gradient (PG)

- On-policy (DQN is off-policy)
- Directly optimize the policy $\pi(s)$

Example: $P(a_1 | s)$ has highest probability




Policy Gradient (PG): Training

- **Steps:**

1. Initialize agent
2. Run policy until termination
3. Record all states, actions and rewards
4. Decrease probability of actions that resulted in low reward
5. Increase probability of actions that resulted in high reward



Example: Pong Game



Policy Gradient (PG): Continuous Action State

- On-policy (DQN is off-policy)
- Directly optimize the policy $\pi(s)$
- **Example:** Assume action space follows Gaussian distribution
- Predict parameters of distribution only
- Sample from distribution based on predicted parameters

Normal distribution, $N(\mu, \sigma^2)$



$$\pi(s) \sim P(a|s) = N(\mu, \sigma^2)$$

Take a sample from distribution = -0.8

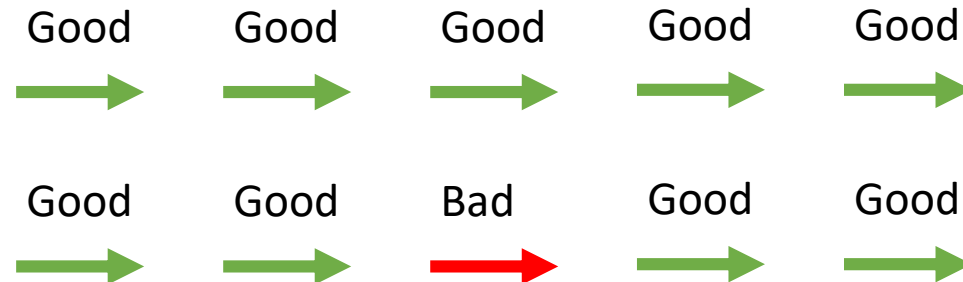
Continuous: How fast left or right should I go?

Discrete: Left or Right?

Policy Gradient (PG): Continuous Action State

- **Weaknesses**

- Needs more data
- Less stable during training
- Poor credit assignment to (s, a) pairs for delayed rewards
 - Calculating reward at the end means all actions will be averaged as good if total reward is high



Questions?